

PETFINDER PAWPULARITY

PREDICTING THE APPEAL OF PET IMAGES



ABHISIKTH PERI

AMULYA KONDA

KASTURI PAL

SANDEEP VENKATA

SHARMISHTA PATRA

Agenda

- Situation
- Complication
- Key Question
- Data structure
- Strategy
- Modeling
- Performance
- Insights

Situation

PetFinder.my, Malaysia's leading animal welfare platform, currently employs a basic Cuteness Meter to rank pet photos. This meter analyzes picture composition and other factors to evaluate the appeal of pet profiles, with over 180,000 animals featured and 54,000 successfully adopted.



Complication

The existing cuteness ranking meter, while helpful, is still in experimental stage and requires significant enhancement to provide precise recommendations.

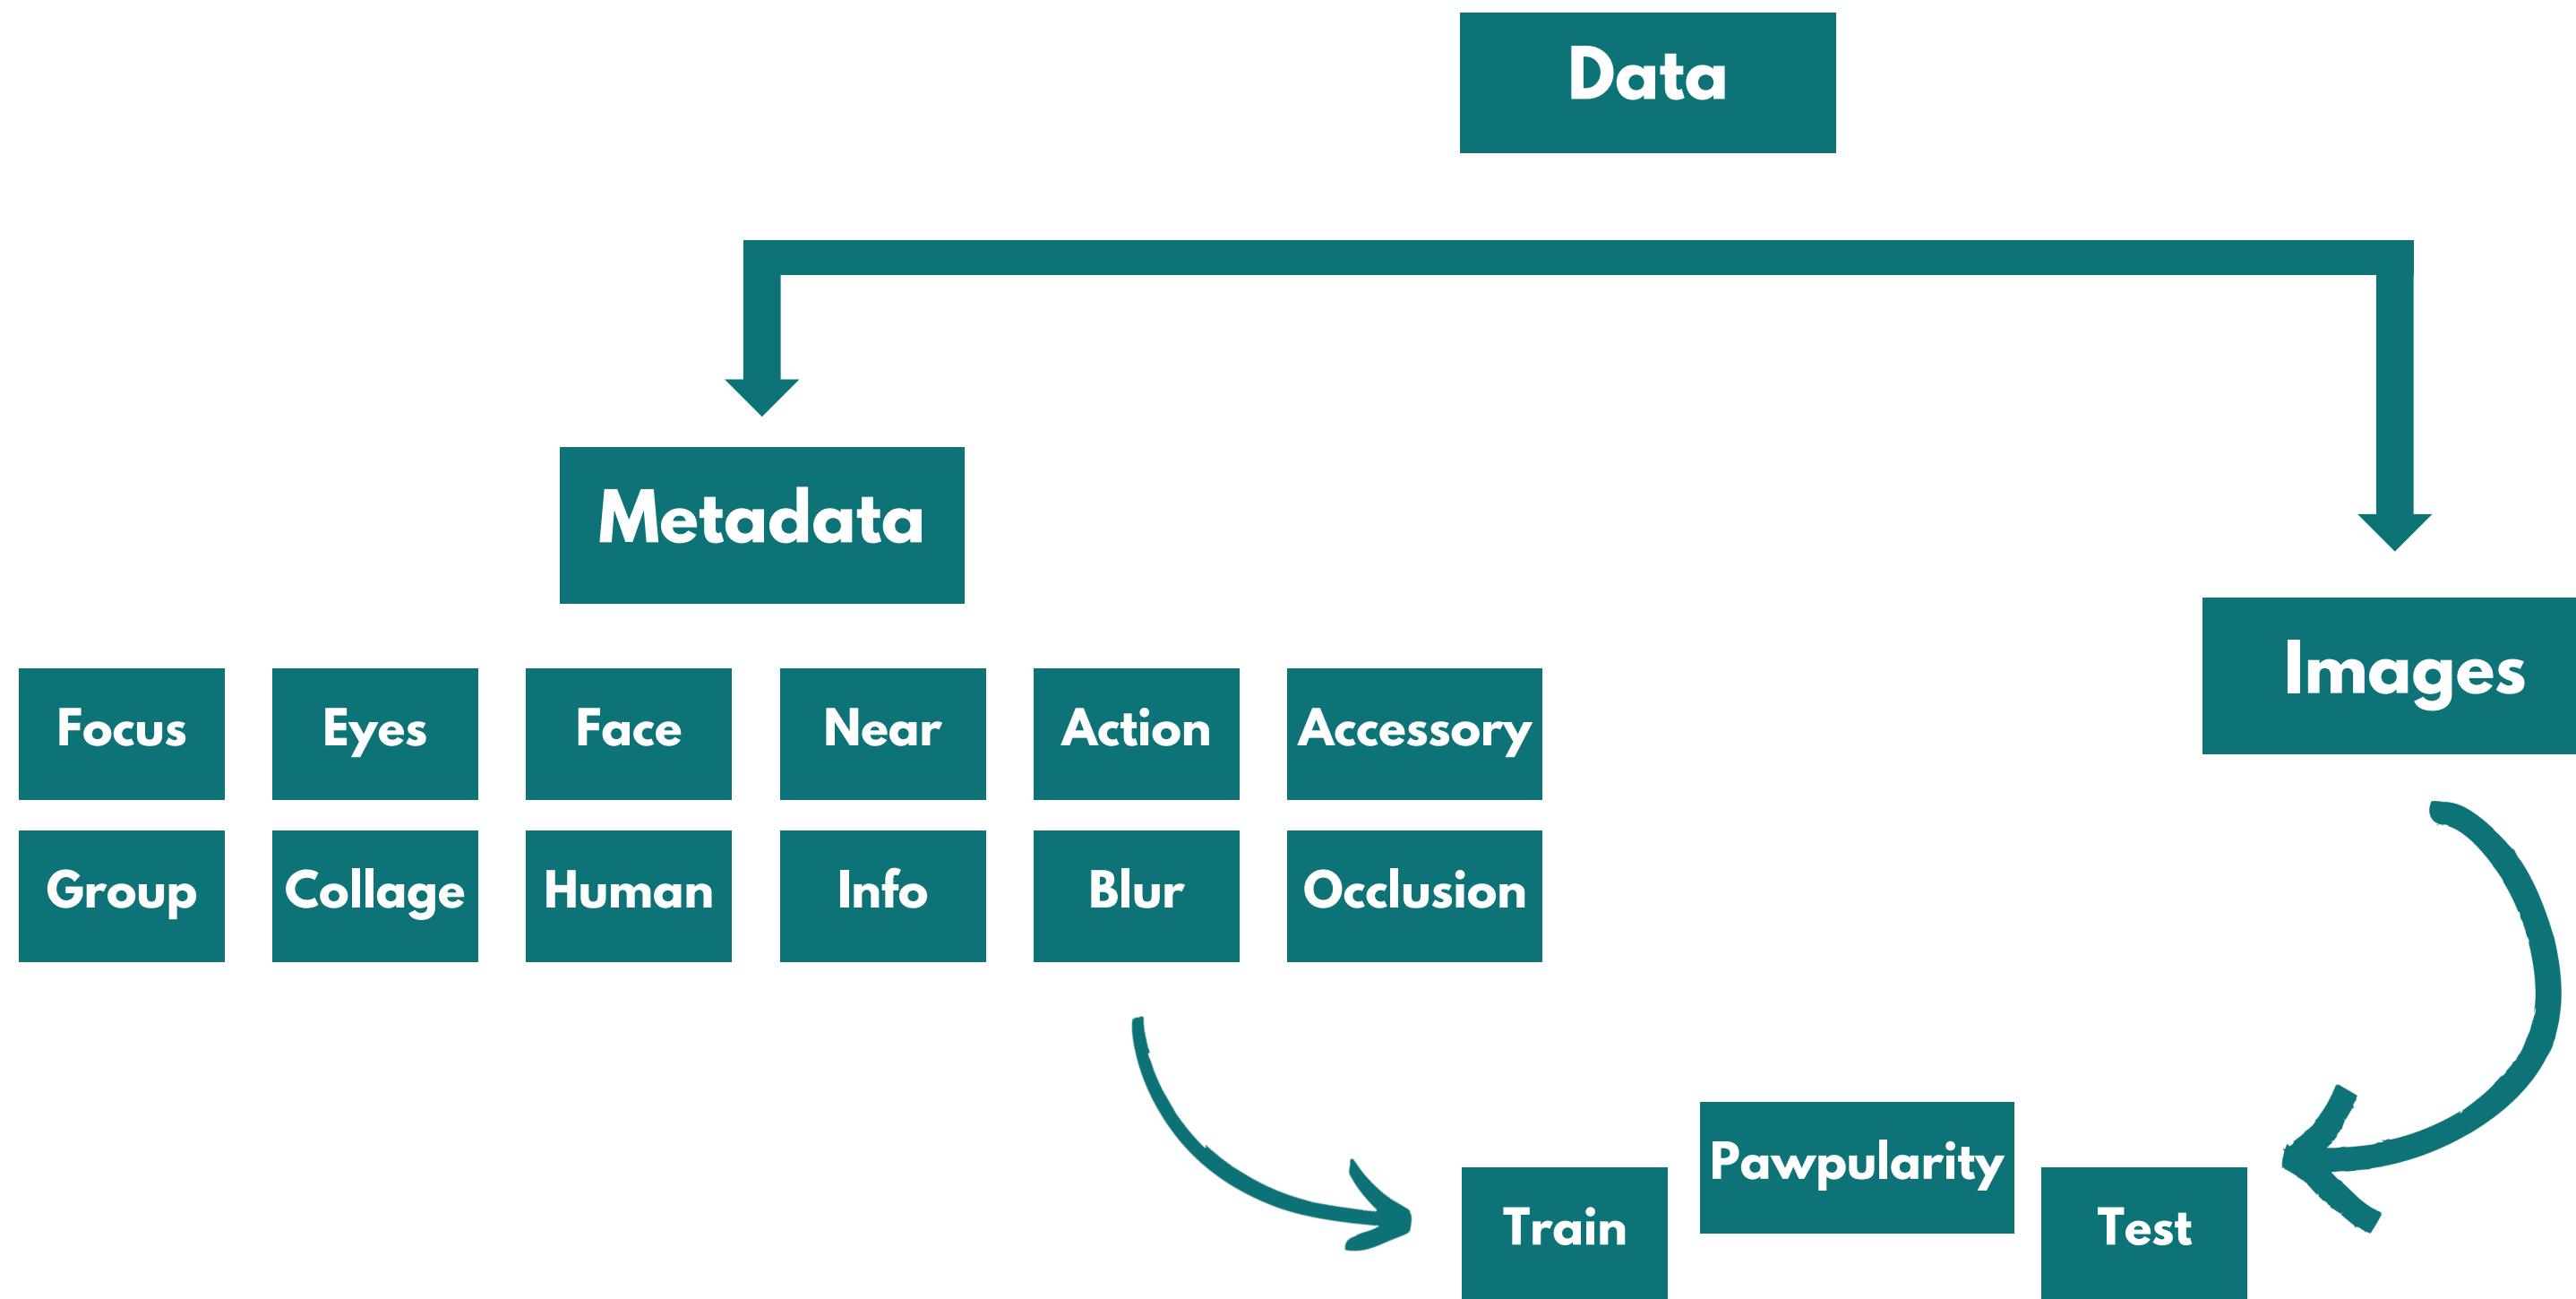


Key Question

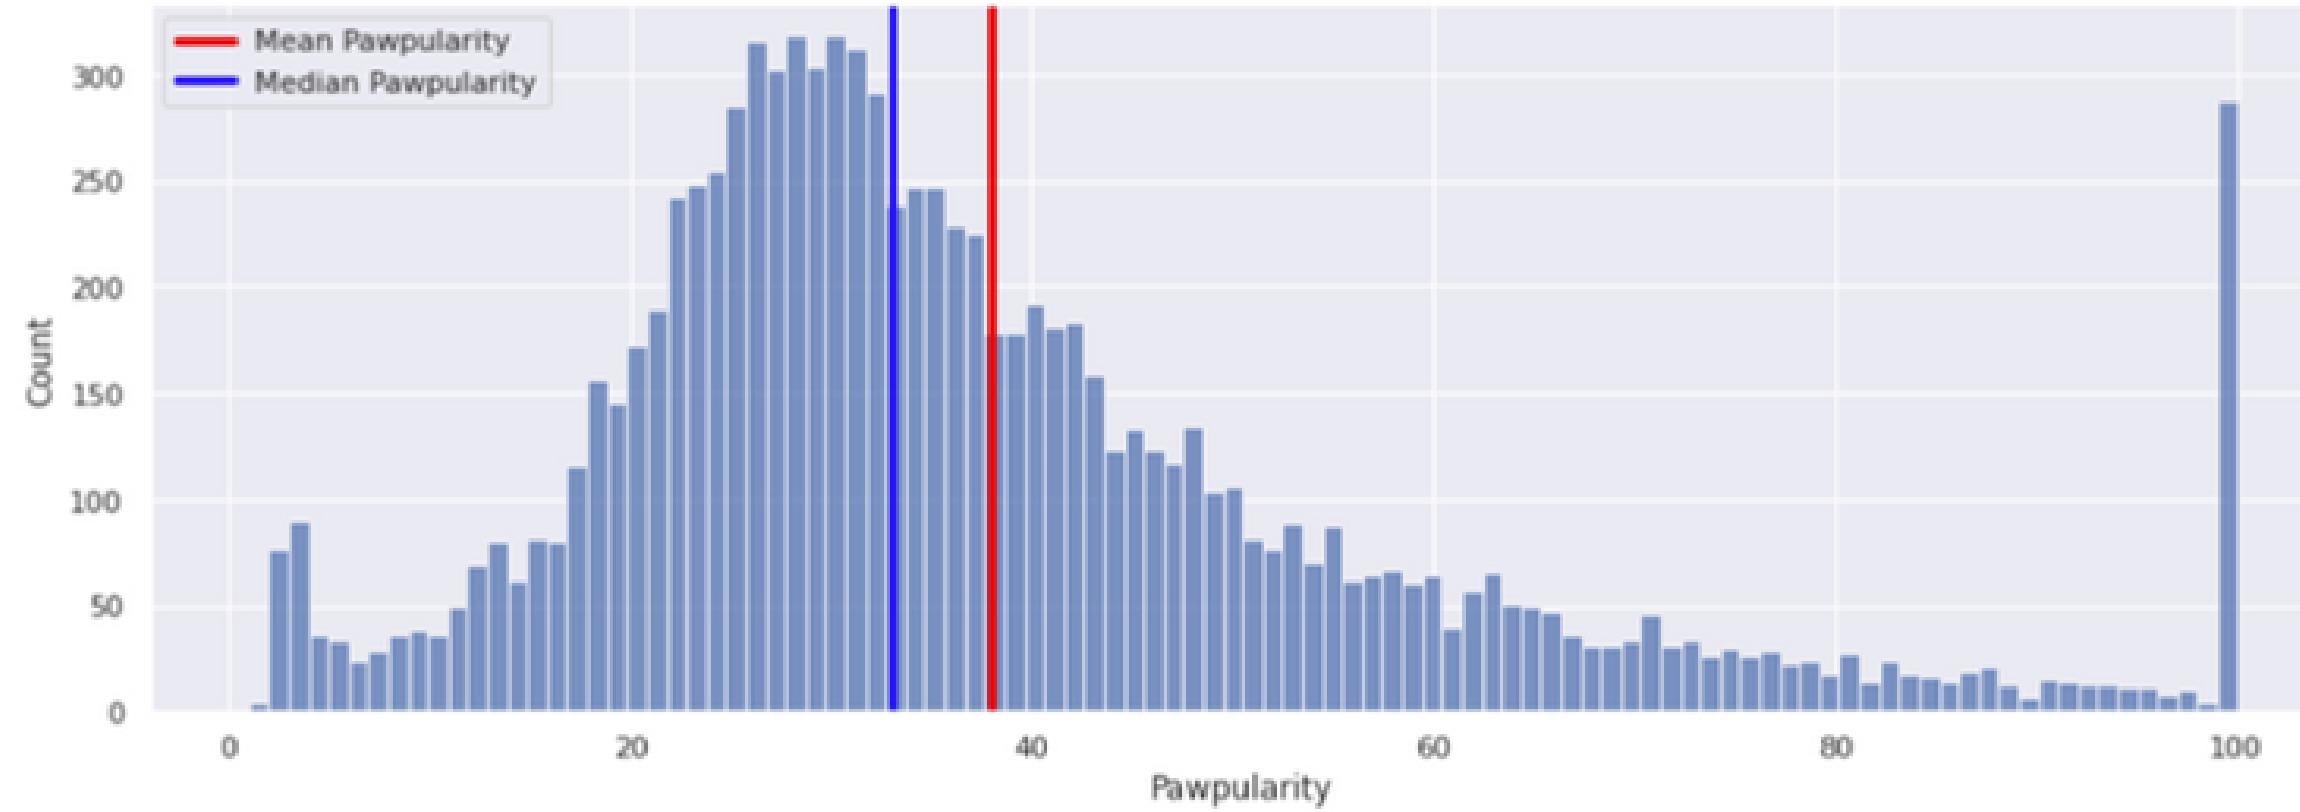
How can we enhance the current prediction algorithm by better analyzing the images to predict the “Pawpularity” of pet photos, ensuring that the model provides accurate recommendations to increase adoption rates?



Data structure



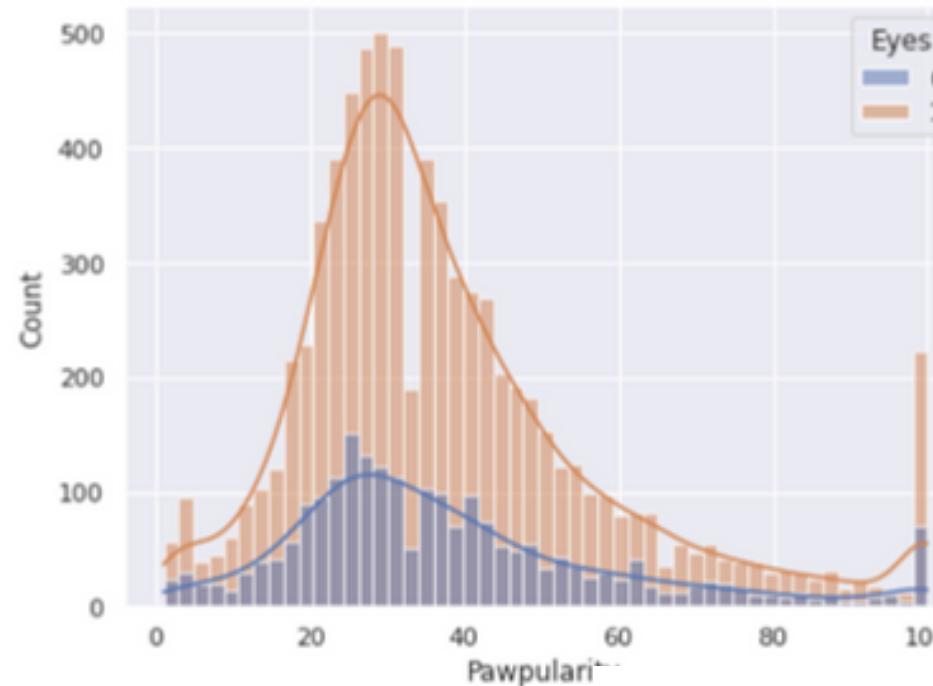
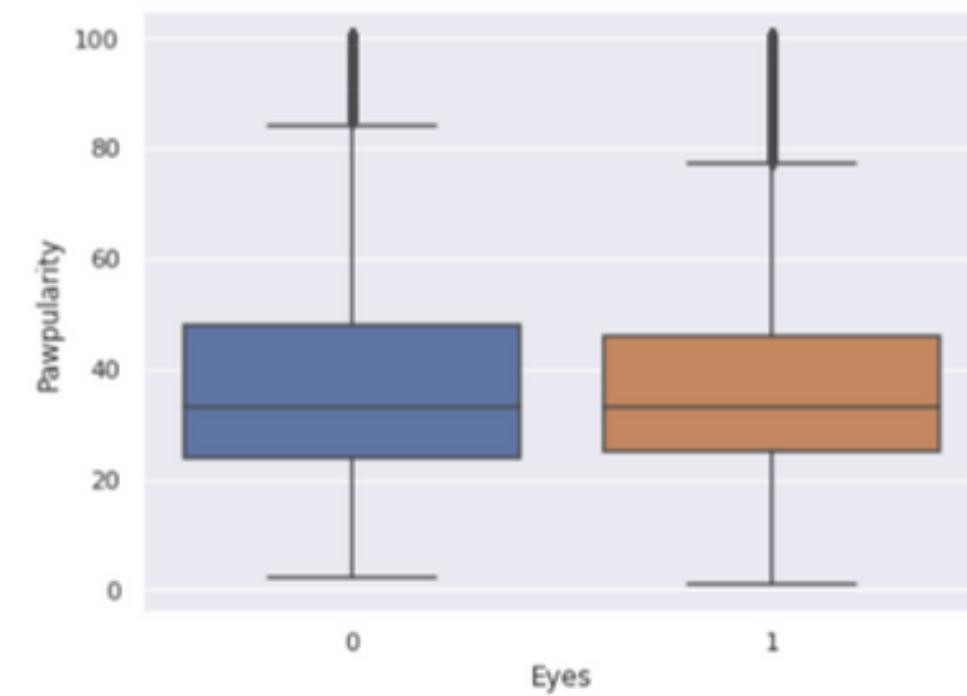
Exploratory Results



- The distribution is **right-skewed**, that is there are a number of images with popularity scores significantly higher than the mode.
- The **median popularity**, or the middle value, is around 30. This means that half of the images have a popularity score below 30 and half have a score above 30.
- There is a **noticeable spike** in the count at around 100 popularity. This could suggest that there are some exceptionally popular images that stand out from the rest

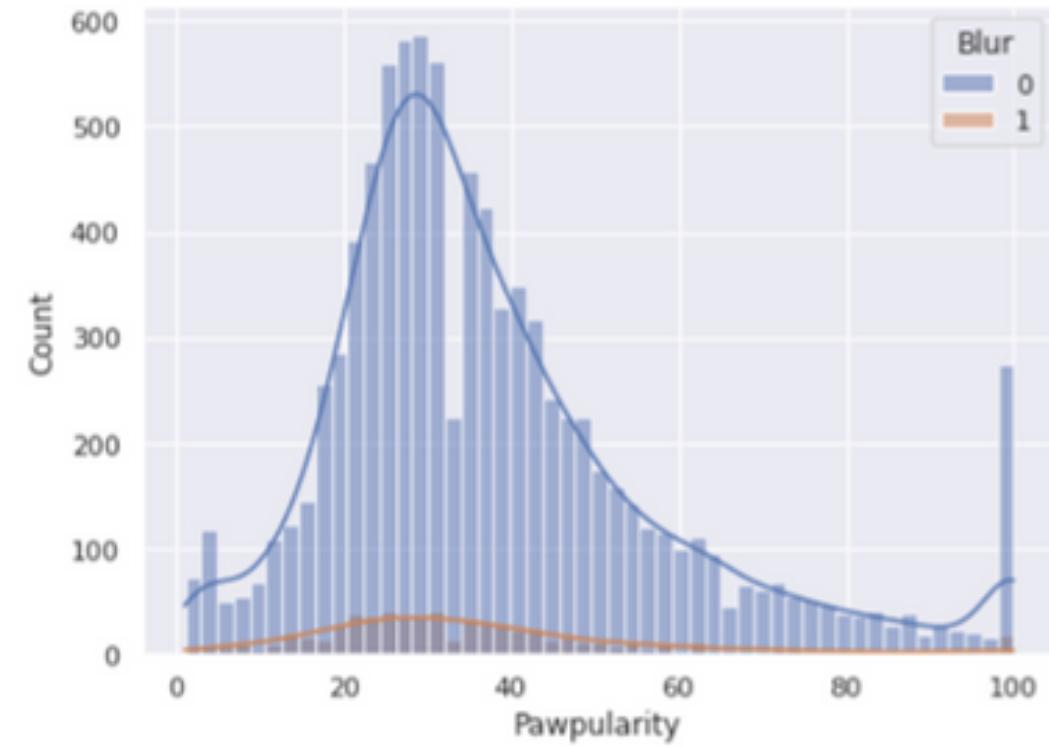
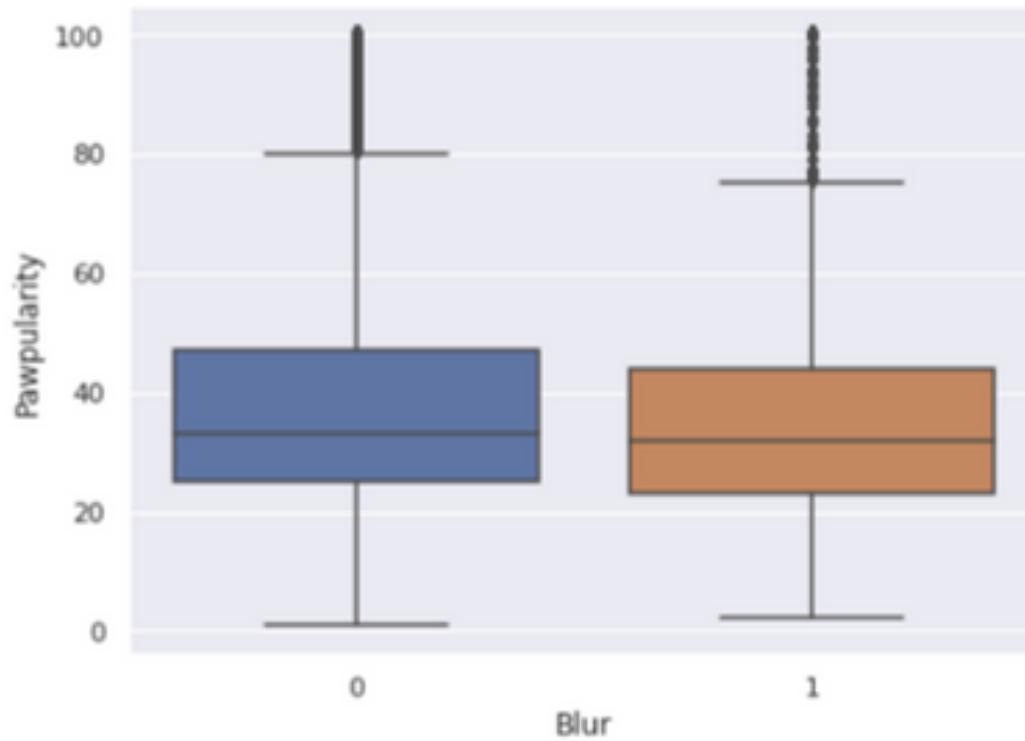
Exploratory Results

Eyes

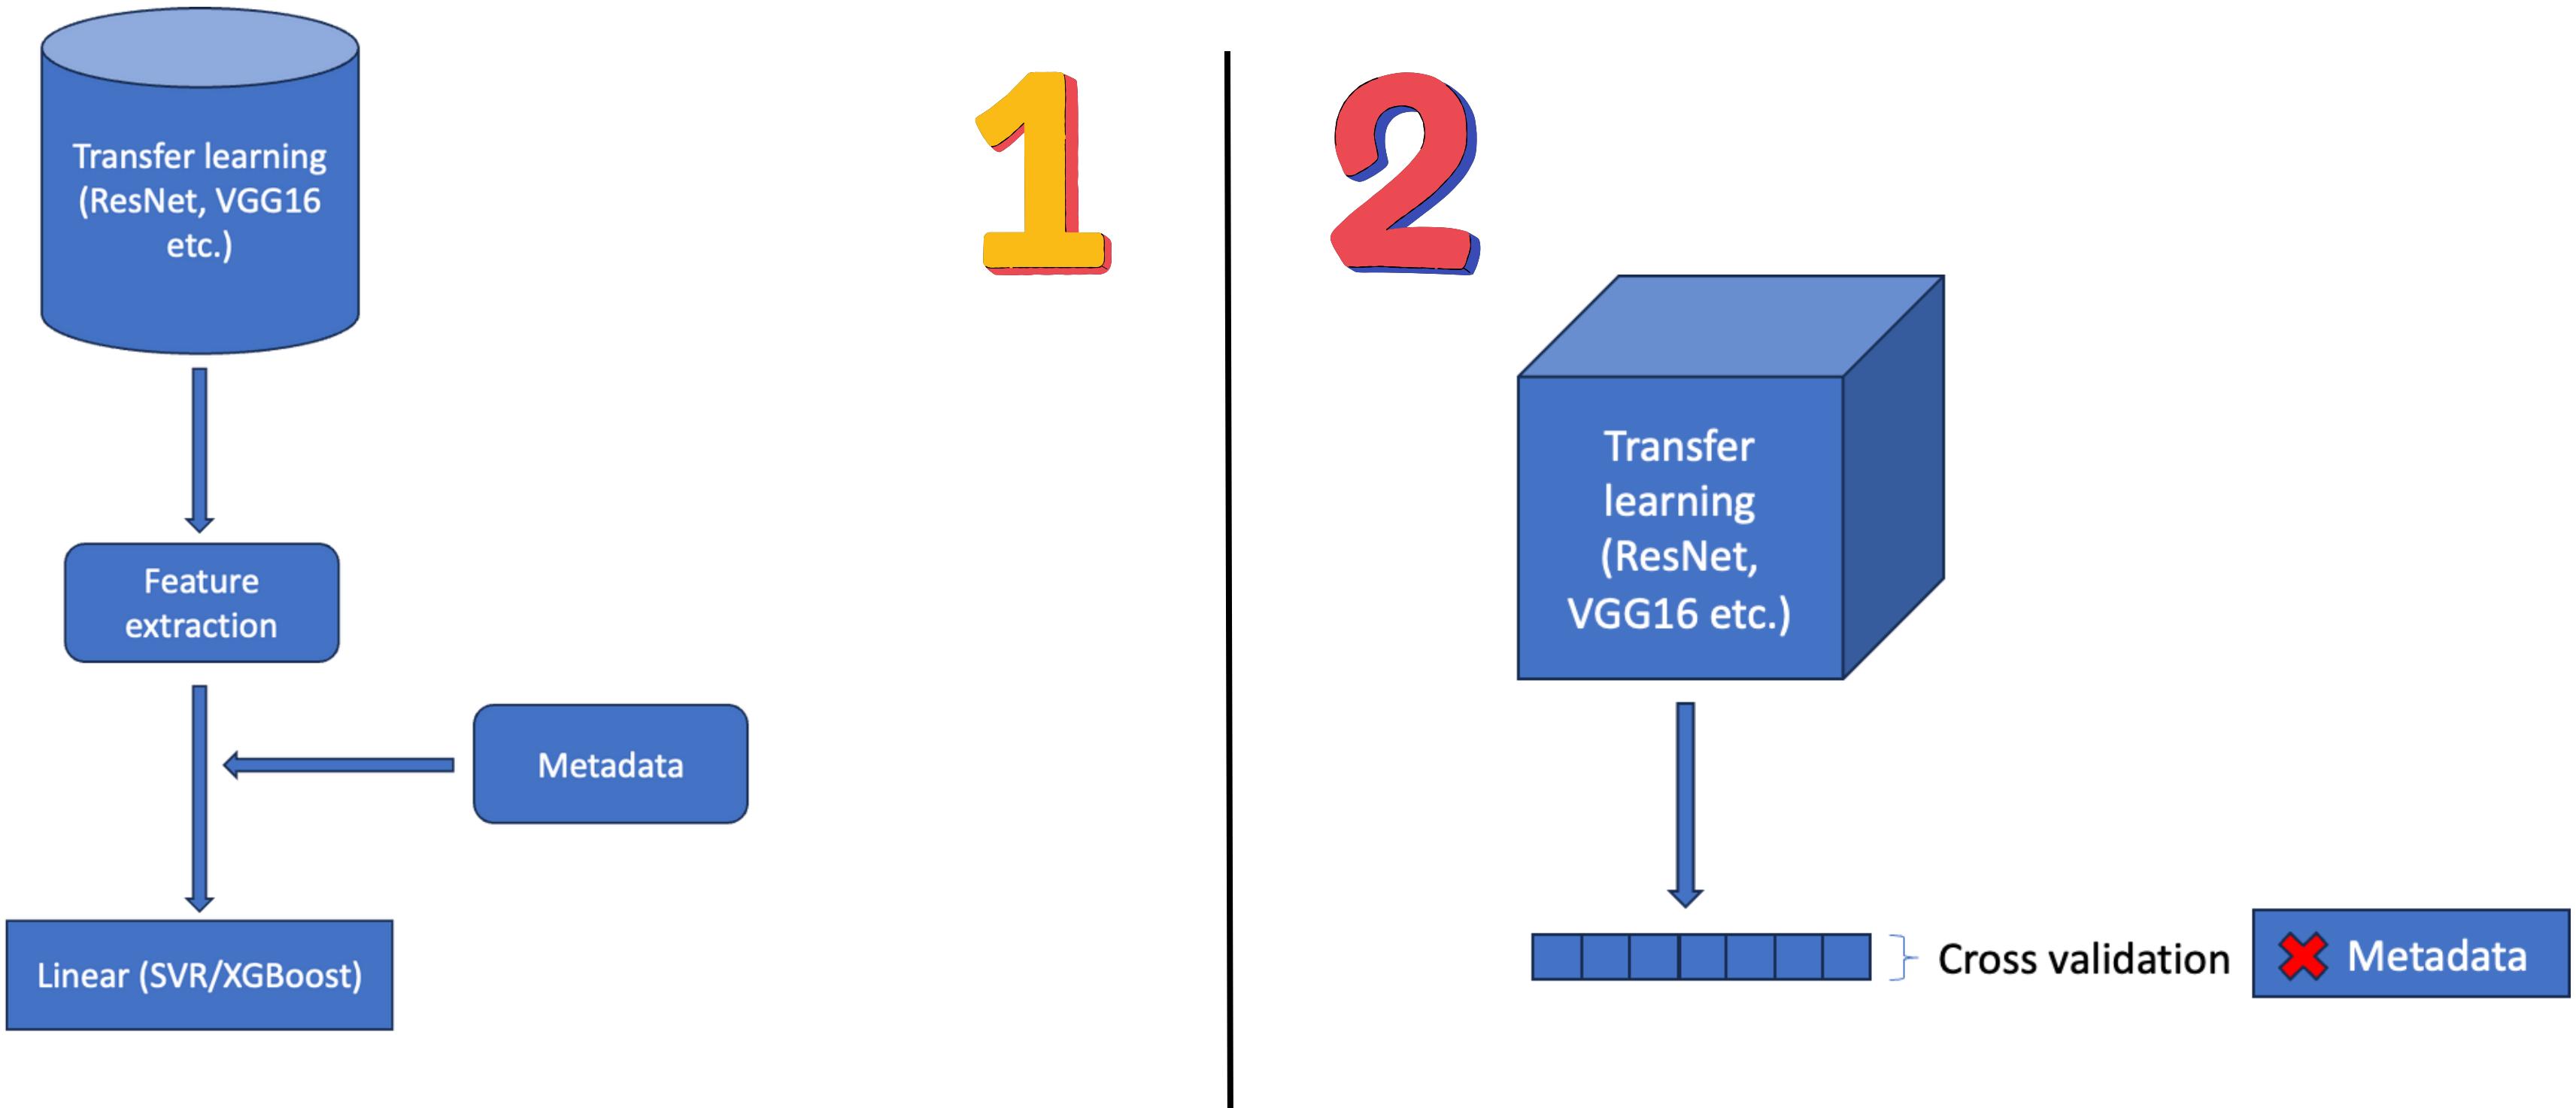


Distribution of Eyes over Pawpularity Score

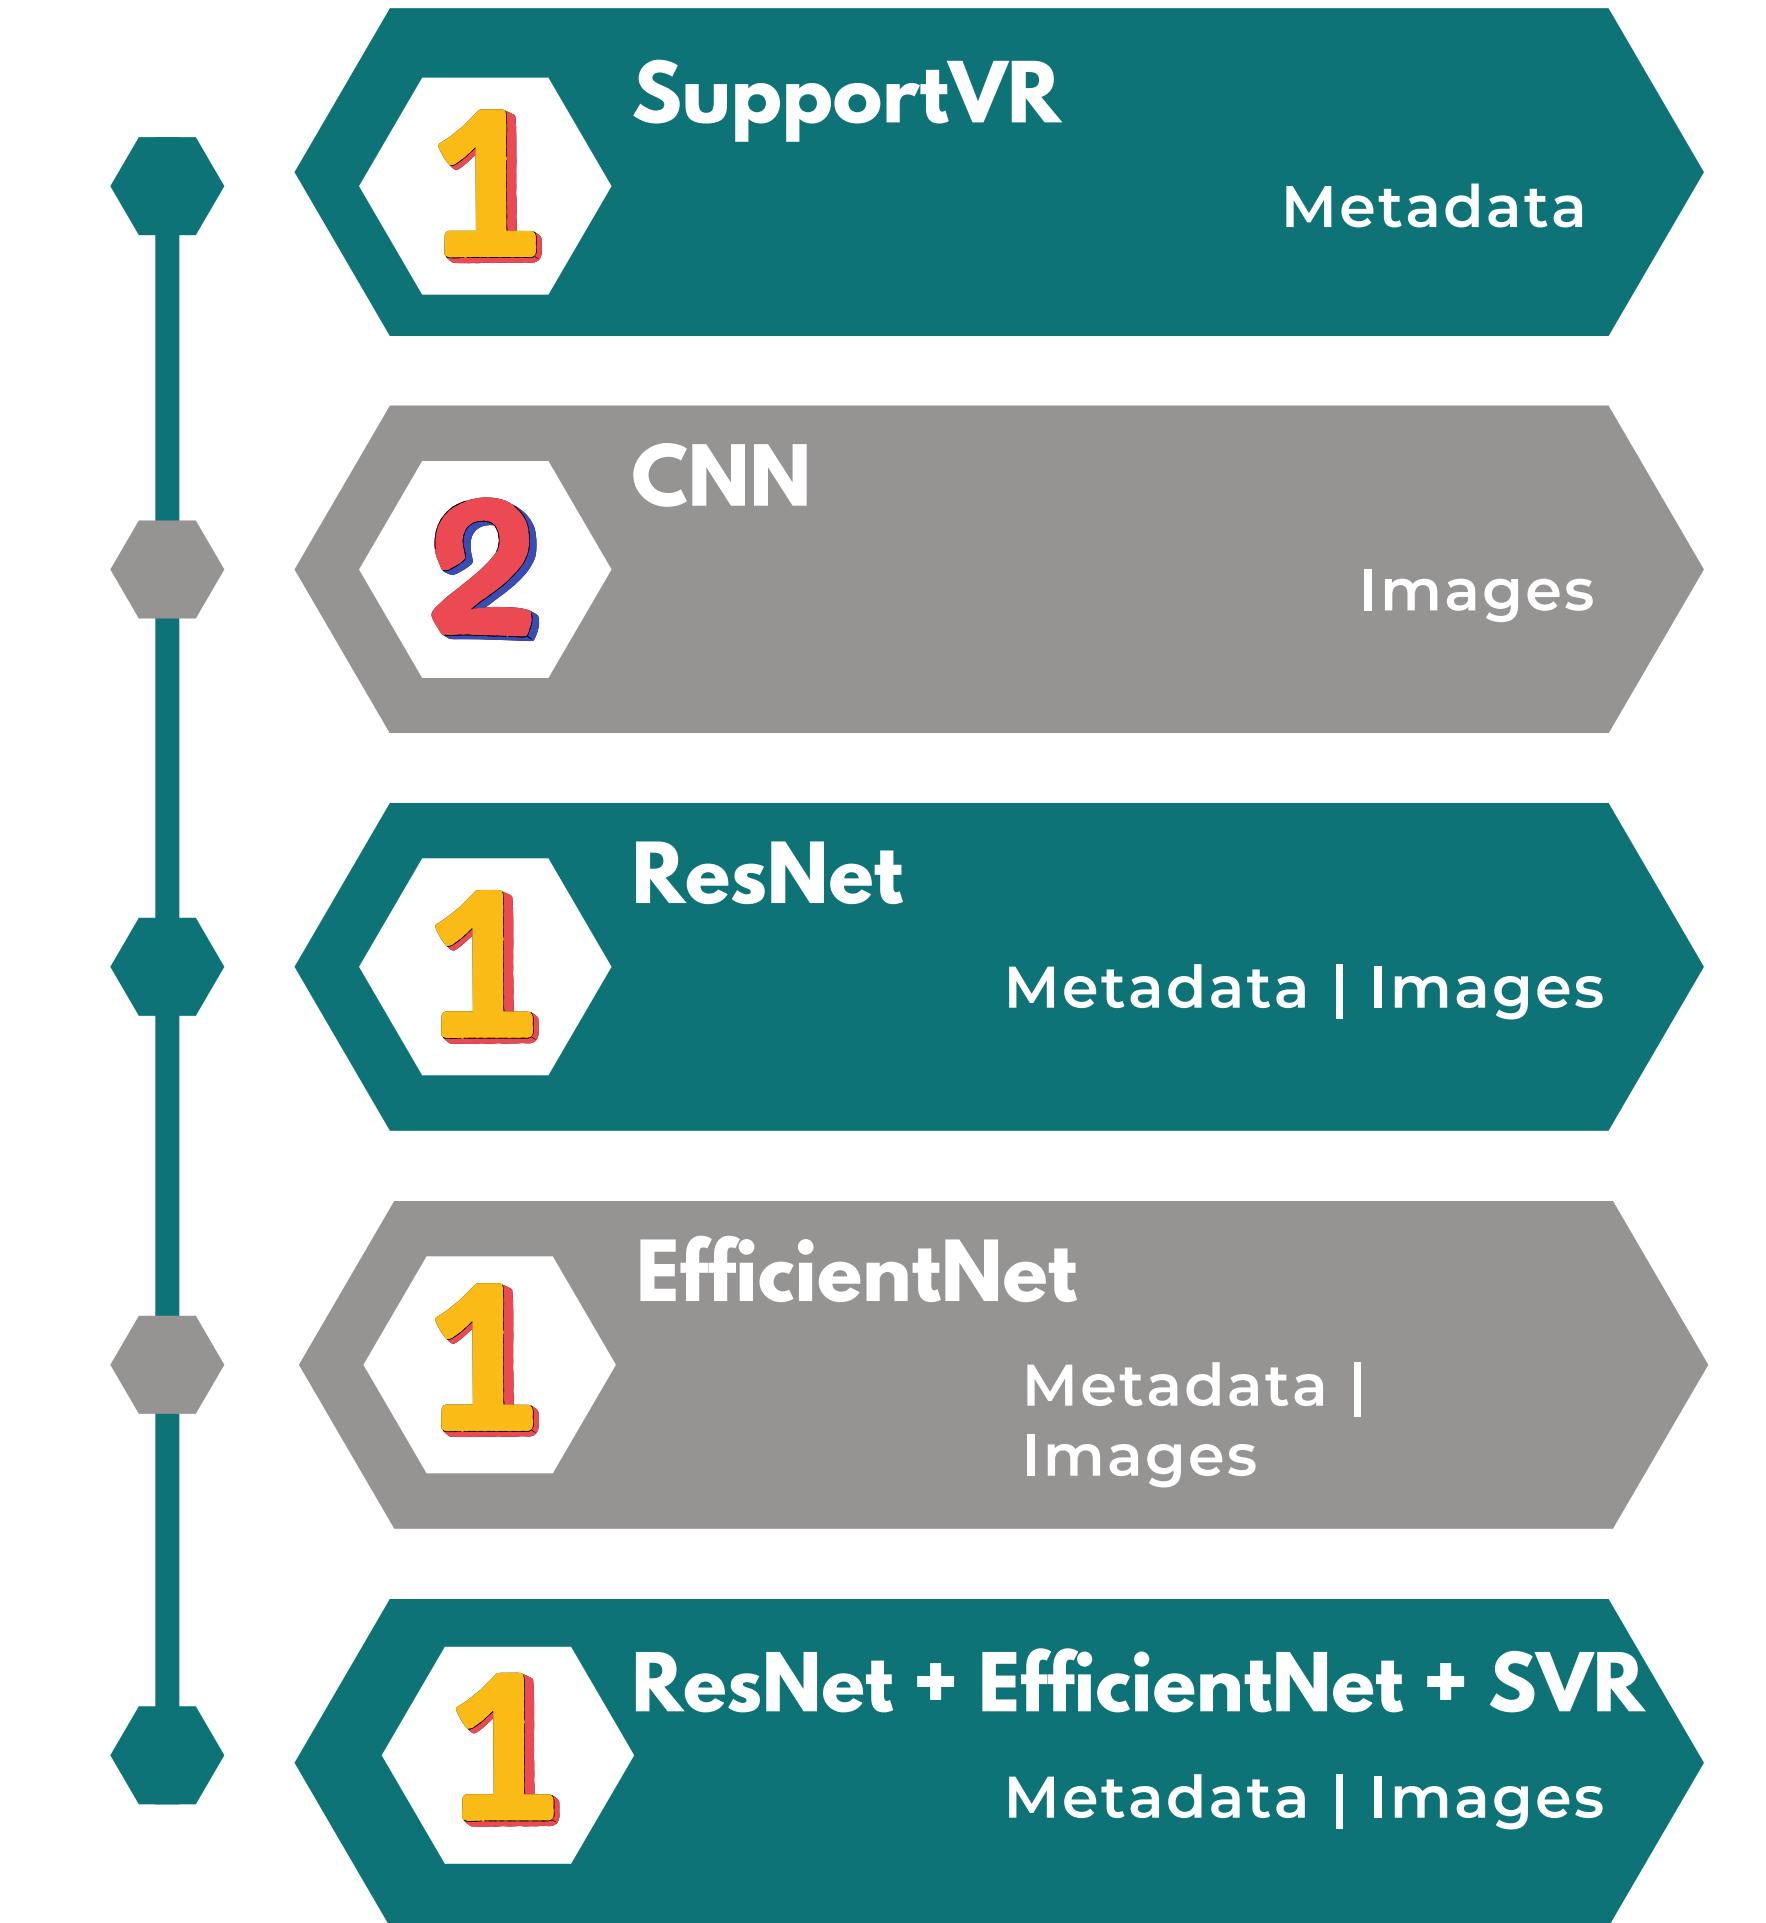
Distribution of Blur over Pawpularity Score



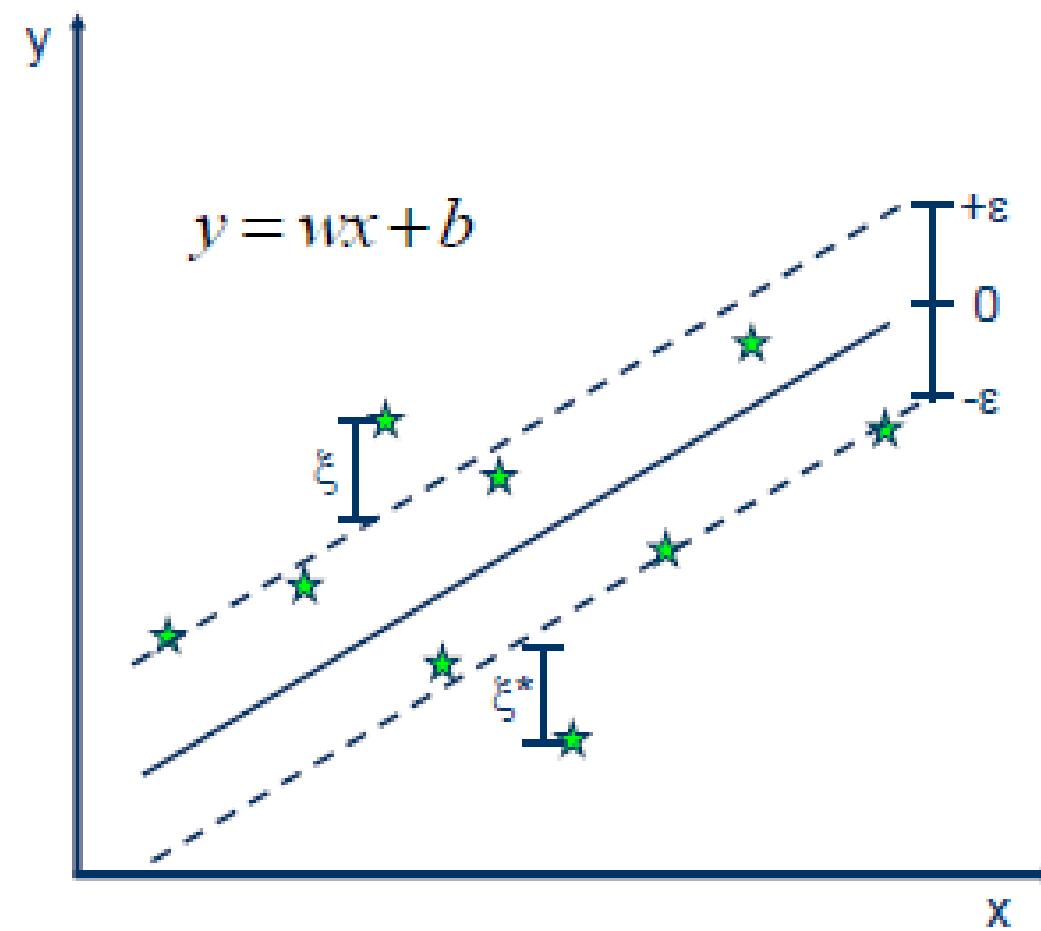
Strategy



Modeling



Support Vector Regression



Linear SVR

- Minimize:

$$\frac{1}{2} \|w\|^2 + C \sum_{i=1}^N (\xi_i + \xi_i^*)$$

- Constraints:

$$y_i - wx_i - b \leq \varepsilon + \xi_i$$

$$wx_i + b - y_i \leq \varepsilon + \xi_i^*$$

$$\xi_i, \xi_i^* \geq 0$$

Model 1: SVR on Metadata

```
Xm = metadata_df.drop(columns=['Id', 'Pawpularity'])
Ym = metadata_df["Pawpularity"]

from sklearn.svm import SVR
from sklearn.metrics import mean_squared_error
SVR = SVR()

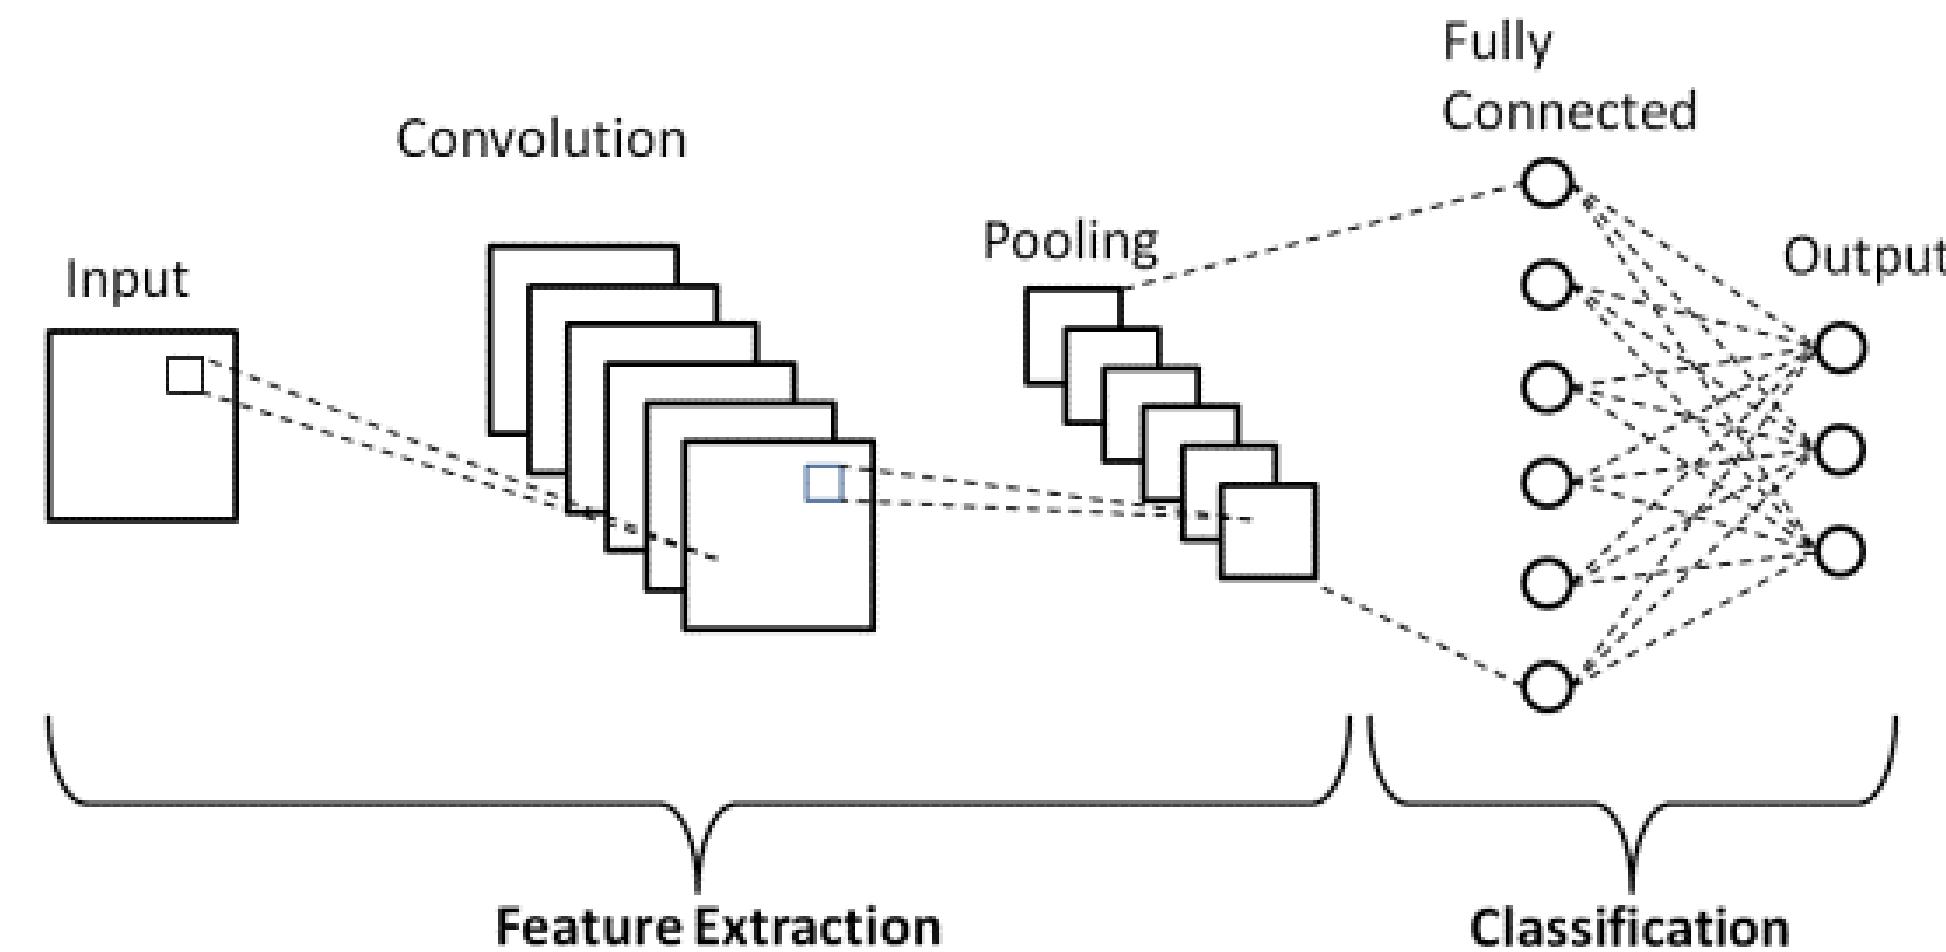
#Fit the model on the training data
SVR.fit(Xm, Ym)

# Make predictions on the testing data
SVR_predictions = SVR.predict(x_test)

print(SVR_predictions)
```

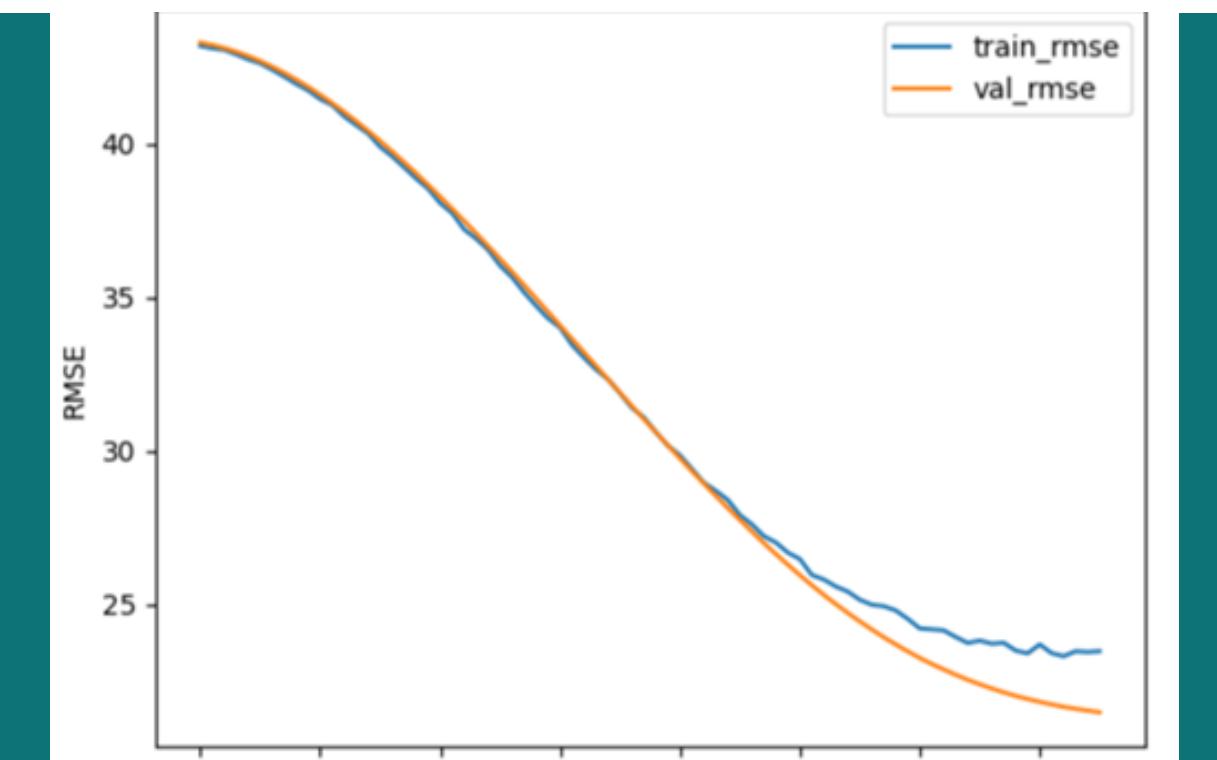
| Submission and Description | Private Score | Public Score | Selected |
|---|---------------|--------------|--------------------------|
|  SVRonmetadata-1 - Version 3 Succeeded (after deadline) · 1m ago · only metadata | 21.08227 | 21.16876 | <input type="checkbox"/> |

CNN architecture



Model 2: CNN on images

```
el = tf.keras.Sequential()  
  
el.add(tf.keras.layers.Conv2D(16,(3,3),activation='relu',input_shape=(128,128,3)))  
el.add(tf.keras.layers.MaxPool2D((2,2)))  
  
el.add(tf.keras.layers.Conv2D(32,(3,3),activation='relu'))  
el.add(tf.keras.layers.MaxPool2D((2,2)))  
  
el.add(tf.keras.layers.Conv2D(64,(3,3),activation='relu'))  
el.add(tf.keras.layers.MaxPool2D((2,2)))  
  
el.add(tf.keras.layers.Flatten())  
el.add(tf.keras.layers.Dense(64,activation='relu'))  
el.add(tf.keras.layers.Dropout(0.1))  
el.add(tf.keras.layers.Dense(1,activation='linear'))  
el.add(tf.keras.layers.Dropout(0.1))  
el.add(tf.keras.layers.Dense(1,activation='linear'))
```



CNN-final - Version 1

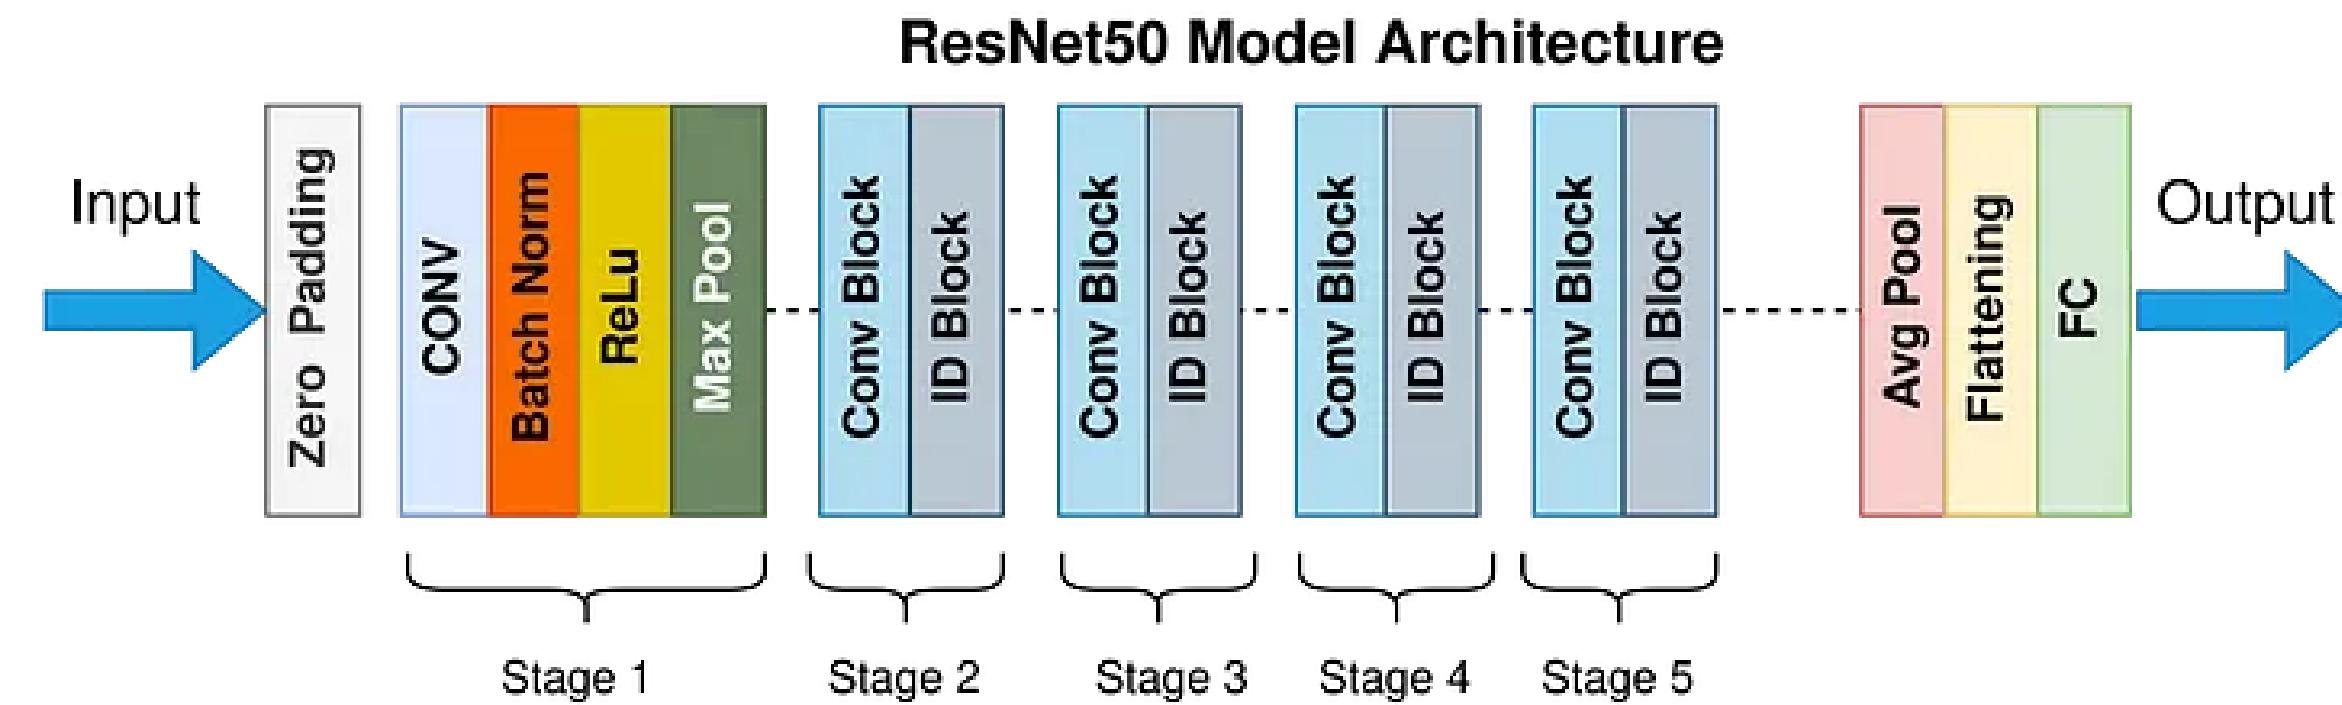
Succeeded (after deadline) · 6h ago · Customized CNN

21.58767

21.65214



Resnet architecture



Model 3: Resnet with metadata

```
tf.keras.backend.clear_session()
models = []
historys = []
kfolds = KFold(n_splits=5, shuffle=True, random_state=42)
train_best_fold = True
best_fold = 0
for index, (train_indices, val_indices) in enumerate(kfolds.split(train)):
    if train_best_fold and index != best_fold: continue
    x_train = train.loc[train_indices, "file_path"]
    tabular_train = train.loc[train_indices, ["Pawpularity"] + columns]
    x_val = train.loc[val_indices, "file_path"]
    tabular_val = train.loc[val_indices, ["Pawpularity"] + columns]
    checkpoint_path = "model_{}_d.h5".format(index)
    checkpoint = tf.keras.callbacks.ModelCheckpoint(checkpoint_path, save_best_only=True)
    early_stop = tf.keras.callbacks.EarlyStopping(min_delta=1e-4, patience=10)
    reduce_lr = tf.keras.callbacks.ReduceLROnPlateau(factor=0.3, patience=2, min_lr=1e-7)
    callbacks = [early_stop, checkpoint, reduce_lr]
    optimizer = tf.keras.optimizers.Adam(1e-3)
    train_ds = tf.data.Dataset.from_tensor_slices((x_train, tabular_train)).map(preprocess).shuffle(512).batch(batch_size).cache().prefetch(2)
    val_ds = tf.data.Dataset.from_tensor_slices((x_val, tabular_val)).map(preprocess).batch(batch_size).cache().prefetch(2)
    model = get_model(image_size, columns)
    model.compile(loss = "mse", optimizer = optimizer, metrics = ["mae", "rmse", "mape"])
    history = model.fit(train_ds, epochs=300, validation_data=val_ds, callbacks=callbacks, batch_size = 8)
    for metric in [("loss", "val_loss"), ("mae", "val_mae"), ("mape", "val_mape"), ("lr")]:
        pd.DataFrame(history.history, columns=metric).plot()
        plt.show()
    model.load_weights(checkpoint_path)
    historys.append(history)
    models.append(model)
```

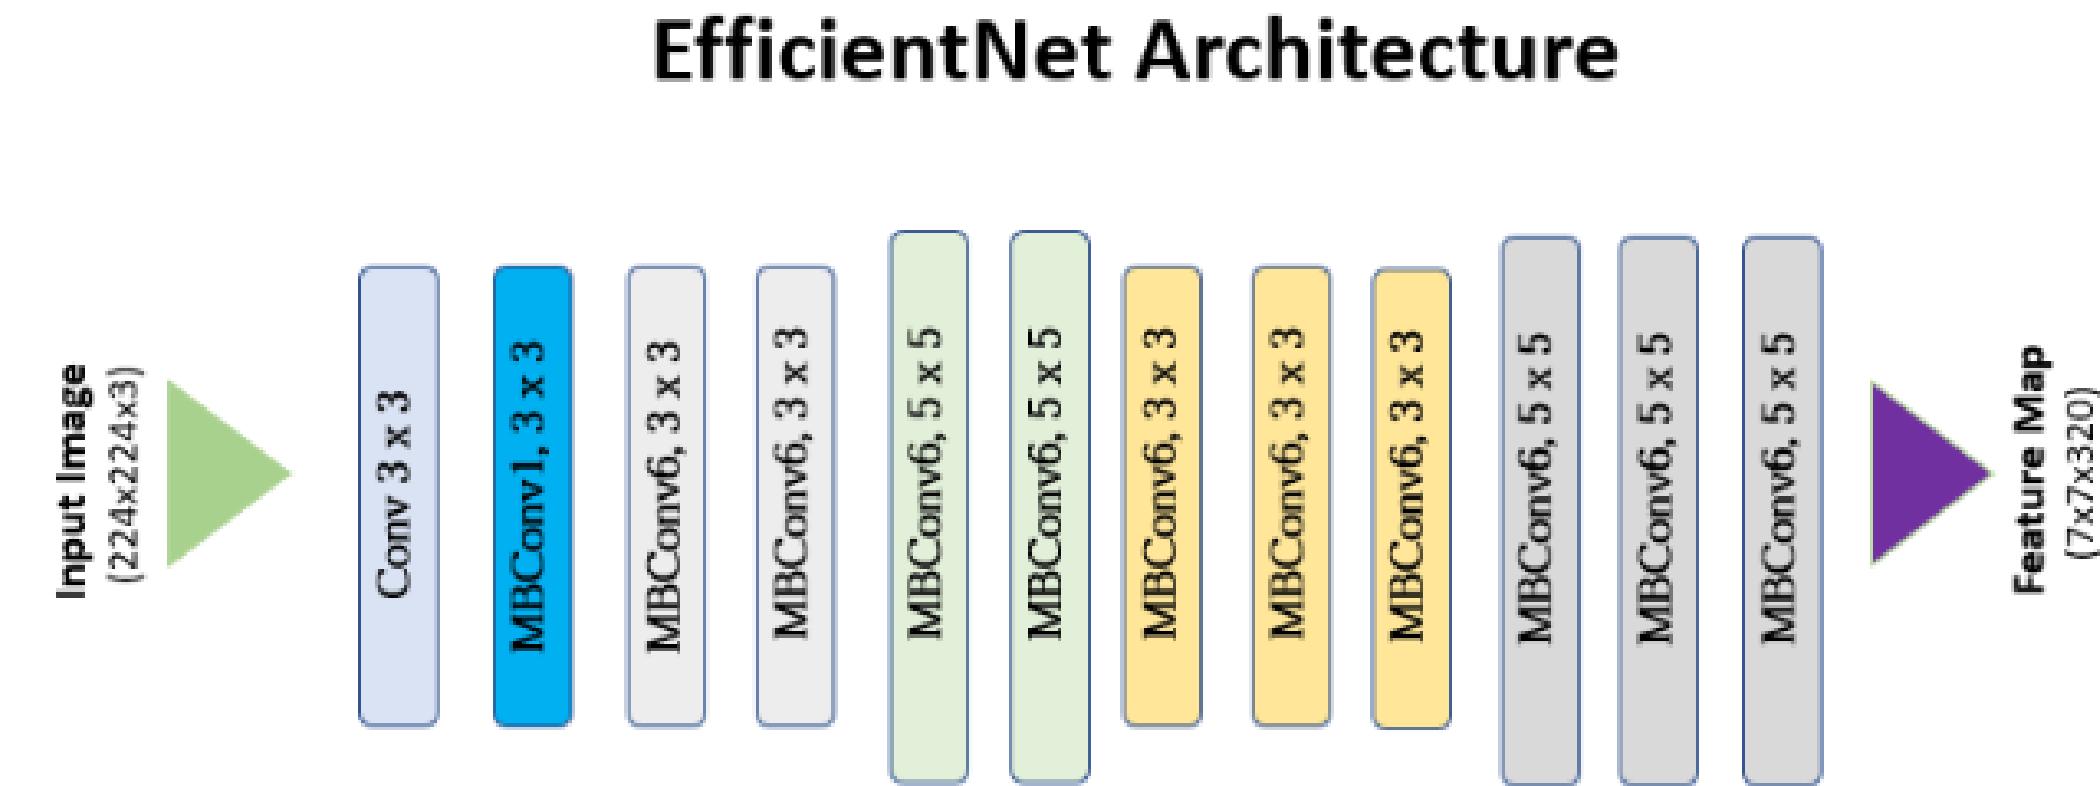
Model 4: Resnet model 3 fine tuned

```
def build_tabular_model(inputs):

    x = keras.layers.Dense(12, activation='relu')(inputs)
    x = keras.layers.Dense(64, activation='relu')(x)
    x = keras.layers.Dropout(0.3)(x)
    x = keras.layers.BatchNormalization()(x)
    x = keras.layers.Dense(128, activation='relu')(x)
    x = keras.layers.Dense(64, activation='relu')(x)
    x = keras.layers.concatenate([x, inputs])
    return x
```

| | | | |
|---|----------|----------|--------------------------|
|  ResnetwithSVR - Version 2 Succeeded (after deadline) · 4d ago | 20.27842 | 20.23979 | <input type="checkbox"/> |
|  Resnetwithmetadata2 - Version 13 Succeeded (after deadline) · 4d ago · resnet with metadata as final layer of neural network | 20.50793 | 20.51378 | <input type="checkbox"/> |

EfficientNet Architecture



Model 5: Efficientnet

```
import os
import tensorflow as tf
from tensorflow.keras.applications import ResNet50
from tensorflow.keras.layers import Dense, GlobalAveragePooling2D, Concatenate, RandomContrast, RandomRotation
from tensorflow.keras.models import Model
from tensorflow.keras.applications import InceptionV3
from tensorflow.keras.applications import EfficientNetB0

def get_model(image_size, columns):
    image_inputs = tf.keras.Input((image_size, image_size, 3))
    tabular_inputs = tf.keras.Input(len(columns))

    #weights_path = '/kaggle/input/resnet50-weights-h5/resnet50_weights_tf_dim_ordering_tf_kernels_notop.h5'
    #if os.path.exists(weights_path):
    #    resnet = ResNet50(include_top=False, weights=weights_path, pooling=None)
    #else:
    #    print("Weights file not found. Using 'imagenet' weights.")
    resnet = tf.keras.applications.EfficientNetB0(include_top=False, weights='imagenet', pooling=None)
    image_x = resnet(RandomContrast(factor=0.1)(RandomRotation(factor=0.15)(image_inputs)))
    image_x = tf.keras.layers.GlobalAveragePooling2D()(image_x)

    # Assuming build_tabular_model is a function you have defined to handle tabular data
    tabular_x = build_tabular_model(tabular_inputs)

    x = tf.keras.layers.Concatenate(axis=1)([image_x, tabular_x])
    output = tf.keras.layers.Dense(1)(x)
    model = tf.keras.Model(inputs=[image_inputs, tabular_inputs], outputs=[output])
    return model
```

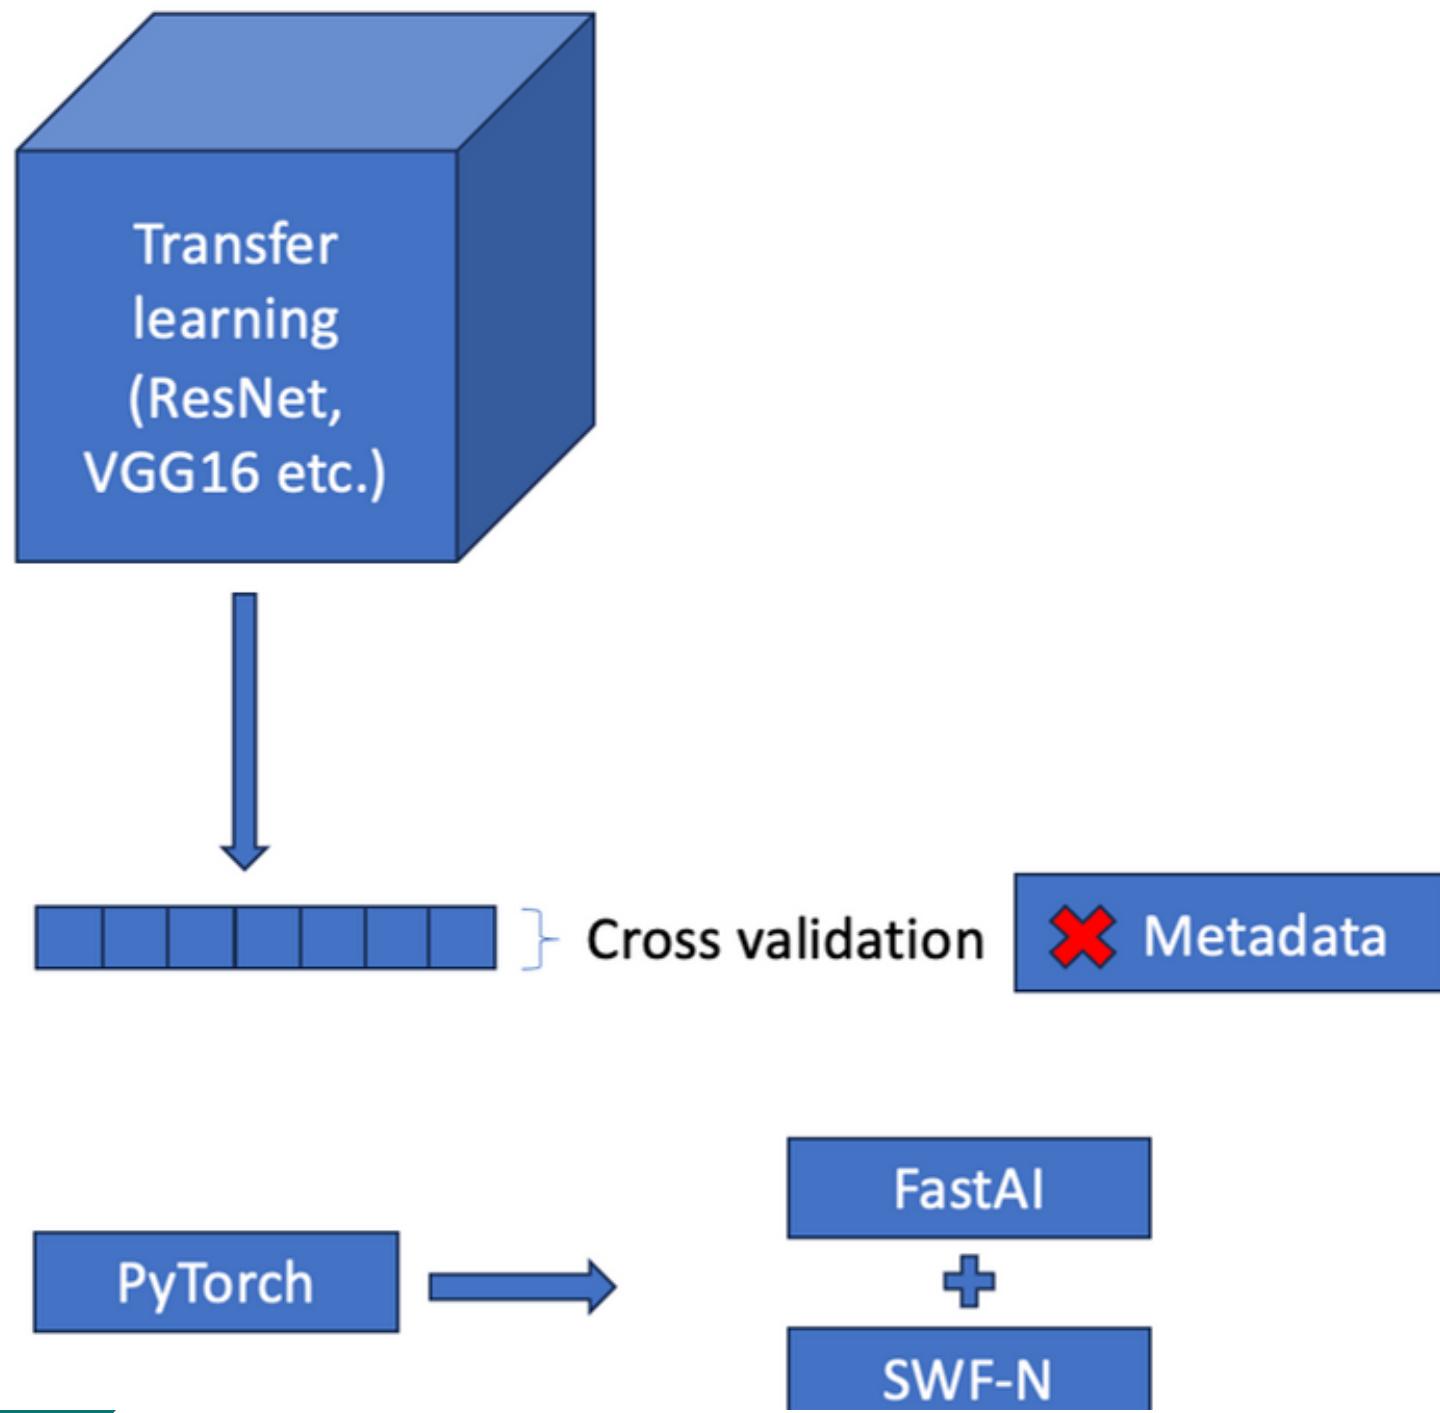
Python · [resnet50_weights.h5, PetFinder.my - Pawpularity Contest](#)

Notebook Input Output Logs Comments (0) Settings

 Competition Notebook Run Best Score
[PetFinder.my - Pawpularity Contest](#) 1317.7s - GPU P100 20.27842 V2

Revised Model approach

Approach 2



Model 6: ResNet+ EfficientNet + SVR

```
# timm_model_lst = ["../input/timm-pretrained-models/beitv2_large_patch16_224.pkl"]
timm_model_lst = ["../input/timm-pretrained-models/tf_efficientnet_l2_ns_475.pkl",
                  "../input/timm-pretrained-models/beitv2_large_patch16_224.pkl"]
#("../input/timm-pretrained-models/swin_base_patch4_window7_224.pkl"
#  "../input/timm-pretrained-models/vit_large_r50_s32_384.pkl",
#  "../input/timm-pretrained-models/vit_large_patch16_384.pkl")
for m in timm_model_lst:
    print("Extracting features using", m)
    pretrained_model = pickle.load(open(m, "rb")).to('cuda')
    trans_config = resolve_data_config({}, model=pretrained_model)
    default_trans = create_transform(**trans_config)

    test_dataset = PetfinderDataSet(img_id_lst=test_id_lst, transform=default_trans, base_path=test_data_base_path)
    test_data_loader = DataLoader(test_dataset, batch_size=batch_size, num_workers=2, shuffle=False)

    with torch.no_grad():
        emb = [pretrained_model(data.to('cuda')).cpu().numpy() for data in test_data_loader]
    EMB_full = np.concatenate(emb, 0)

    model_name = m.split("/")[-1].split(".")[0]
    pickle.dump(EMB_full, open("./testset_features/" + model_name + "_features.pkl", "wb"))

del pretrained_model, EMB_full
torch.cuda.empty_cache()
gc.collect()
print("Done!\n")
```

Model 6: ResNet+EfficietNet + SVR

```
# Fixed setting
models = ["RN50x16",
          "RN50x4",
          "ViT-B-16",
          "ViT-B-32",
          "tf_efficientnet_l2_ns_475"]
model_old_hill_1 = ["RN50x16",
                     "beitv2_large_patch16_224",
                     "beitv2_large_patch16_224"] # * 1.032
model_old_hill_2 = ["ViT-B-16",
                     "beitv2_large_patch16_224"] # * 1.032
models_man_2 = ["RN50x16",
                "RN50x4",
                "ViT-B-16",
                "ViT-B-32",
                "tf_efficientnet_l2_ns_475",
                "beitv2_large_patch16_224"] # loss: 17.296 * 1.030
models_man_5 = ["RN50x16",
                "RN50x4",
                "ViT-B-16",
                "ViT-B-32",
                "beitv2_large_patch16_224",
                "swin_base_patch4_window7_224",
                "tf_efficientnet_l2_ns"] # loss: 17.257 * 1.029
# "beitv2_large_patch16_224"
#     "vit_large_patch16_384",
#     "vit_large_r50_s32_384"
# models = ["beitv2_large_patch16_224",
#           "RN50x16"]
models = models_man_2
EMB_test = None
for i in range(len(models)):
    if i == 0:
        EMB_test = pickle.load( open( './testset_features/' + models[0] + '_features.pkl', "rb" ) )
    else:
        EMB_test = np.concatenate((EMB_test, pickle.load( open( './testset_features/' + models[i] + '_features.pkl', "rb" ) )), axis=1)
```

Submission and Description

Private Score



ResnetandefficientSVR - Version 2
Succeeded (after deadline) · 21h ago

17.0998

Performance

| Model | RMSE |
|----------------------------|-------|
| SVR | 21.08 |
| CNN | 21.59 |
| ResNet | 20.28 |
| EfficientNet | 20.28 |
| ResNet + EfficientNet+ SVR | 17.08 |

Insights

- Positive Contribution of Metadata
- CNN Limitations
- Transfer Learning Superiority
- Multiple Transfer Learning Models
- Effective Hyperparameter Tuning
- Significance of Data Augmentation





**THANK
YOU**